

Name: MD Ekram Ullah
Batch: ML-DS 102

1.Problem Statement:

Given the details of customers in a bank, the aim is to improve bank marketing strategies by analyzing their past marketing strategies and recommending which customers to target.

2.Approach:

First of all, I loaded the data in jupyter notebook and put a detailed look into the data. I performed EDA vastly to look into categorical and numerical feature distribution, relationship between categorical features and label, relationship between discrete numerical features and label and also continuous numerical features and label, pair plot and so on. After getting a clear idea of my data, I performed Feature engineering into the dataset. Where I dropped the unwanted features, handled the missing values, handles categorical features, feature scaling and removed the outliers. Then after getting a clean dataset, I splitted the dataset into training and test set and then I trained the models.

3.Phases:

i) Exploratory Data Analysis:

I did my EDA task vastly with so many steps in it which are described below:

a) Find Missing Values:

Here I tried to find the missing values feature by feature. But found no missing values.

2. Find Missing Values

```
In [11]: # find missing values
features_na = [features for features in df.columns if df[features].isnull().sum() > 0]
for feature in features_na:
    print(feature, np.round(df[feature].isnull().mean(), 4), ' % missing values')
else:
    print("No missing value found")

No missing value found
```

b) Explore the Categorical Features:

Here I looked into details of categorical features.

4. Explore the Categorical Features

```
In [13]: categorical_features=[feature for feature in df.columns if ((df[feature].dtypes=='O') & (feature not in categorical_features))
categorical_features
```

```
Out[13]: ['job',
          'marital',
          'education',
          'default',
          'housing',
          'loan',
          'contact',
          'month',
          'poutcome']
```

```
In [14]: for feature in categorical_features:
          print('The feature is {} and number of categories are {}'.format(feature,df[feature].nunique()))

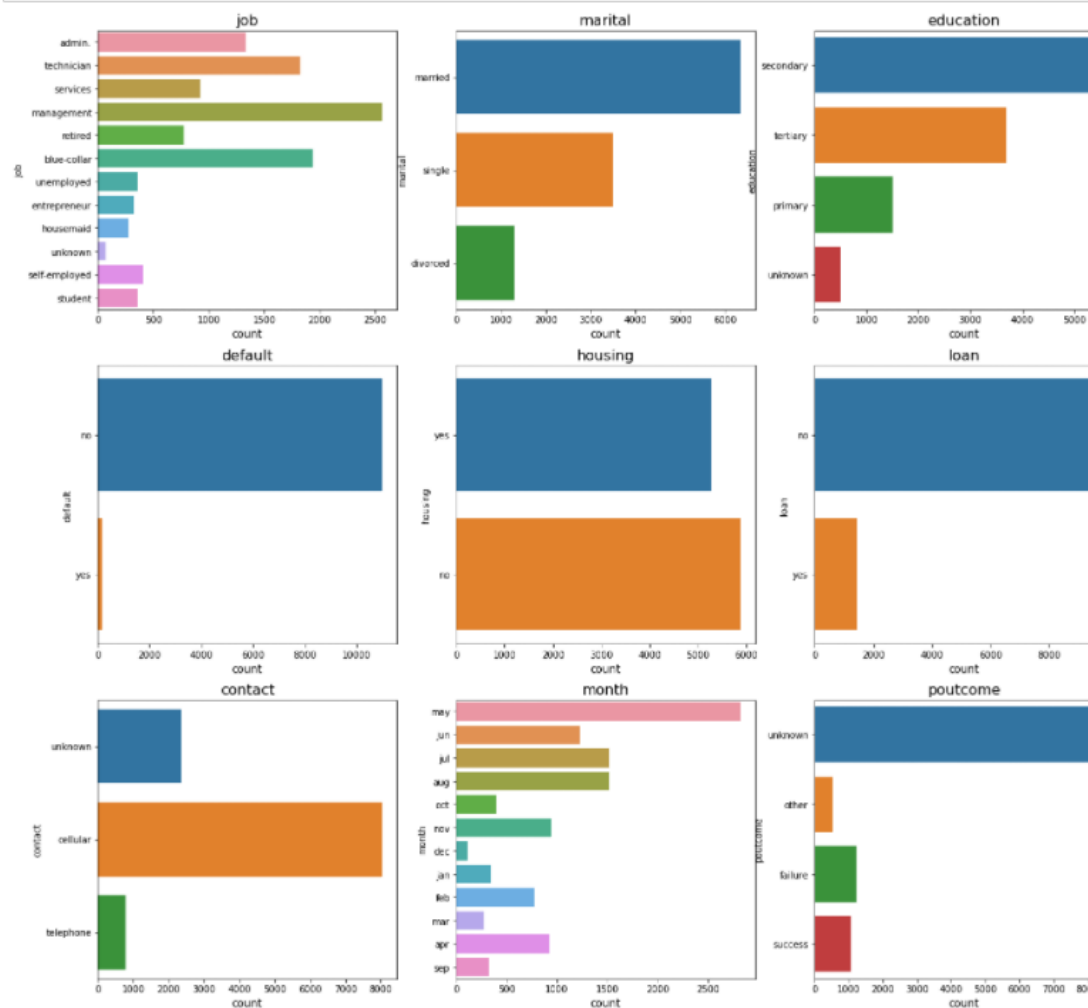
The feature is job and number of categories are 12
The feature is marital and number of categories are 3
The feature is education and number of categories are 4
The feature is default and number of categories are 2
The feature is housing and number of categories are 2
The feature is loan and number of categories are 2
The feature is contact and number of categories are 3
The feature is month and number of categories are 12
The feature is poutcome and number of categories are 4
```

c) Find Categorical Feature Distribution:

I looked into the distribution of every categorical feature, how there unique values are distributed and the frequency.

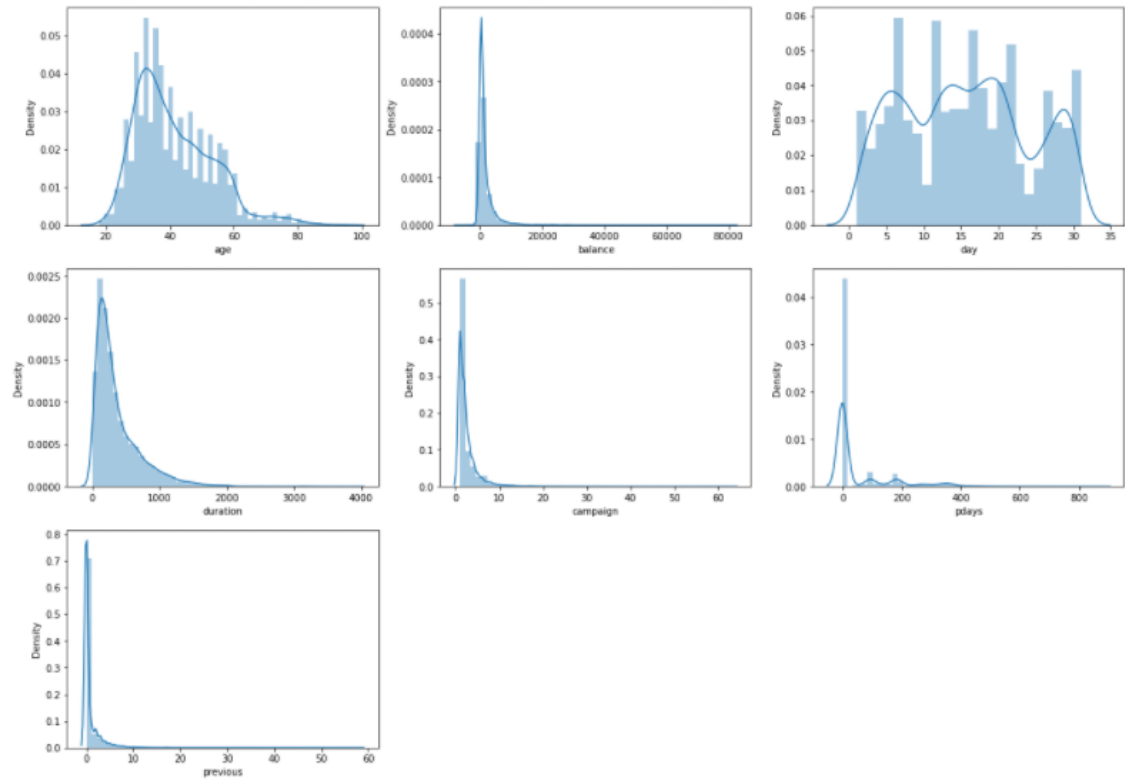
```
In [15]: #check count based on categorical features
plt.figure(figsize=(20,80), facecolor='white')
plotnumber = 1
for categorical_feature in categorical_features:
    ax = plt.subplot(12,3,plotnumber)
    plt.title(categorical_feature,fontsize=16)
    plt.xlabel('frequency',fontsize=12)

    sns.countplot(y=categorical_feature,data=df)
    plotnumber+=1
plt.show()
```



d) **Distribution of Continuous Numerical Features:**

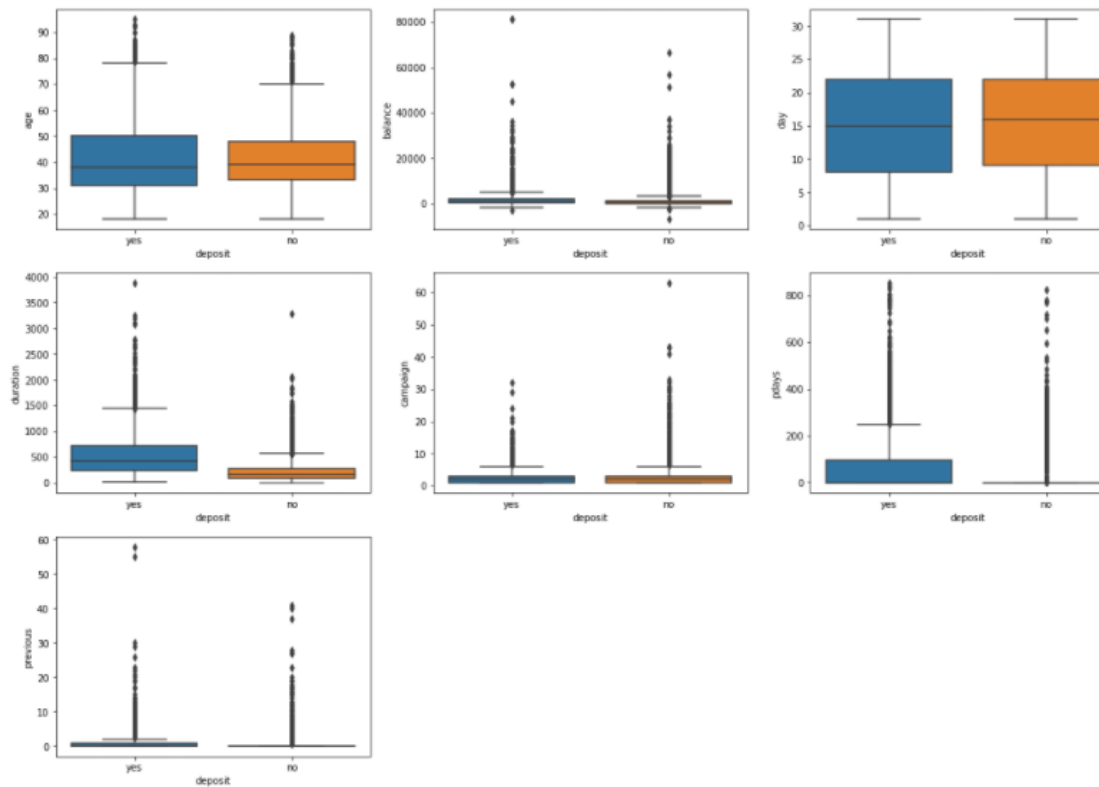
```
In [21]: #plot a univariate distribution of continues observations
plt.figure(figsize=(20,60), facecolor='white')
plotnumber =1
for continuous_feature in continuous_features:
    ax = plt.subplot(12,3,plotnumber)
    sns.distplot(df[continuous_feature])
    plt.xlabel(continuous_feature)
    plotnumber+=1
plt.show()
```



e) **Relation between Continous numerical Features and Labels:**

I tried to look the relation between each feature and label.

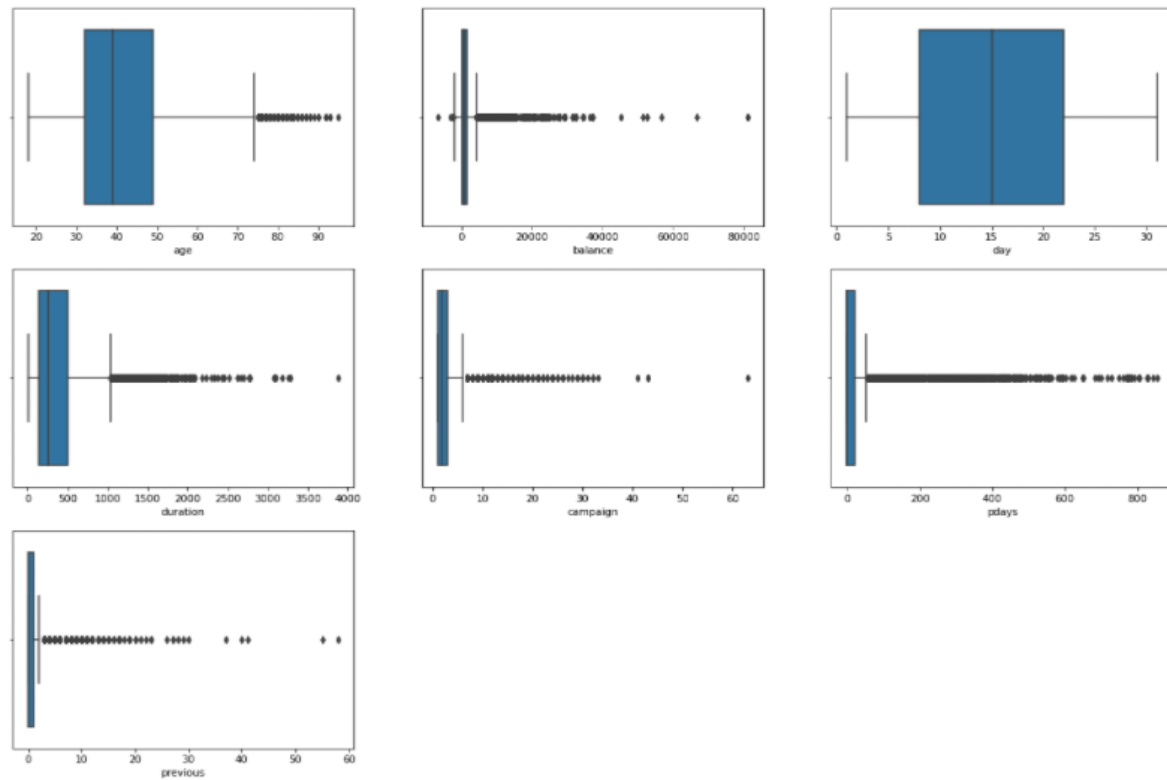
```
In [22]: #boxplot to show target distribution with respect numerical features
plt.figure(figsize=(20,60), facecolor='white')
plotnumber = 1
for feature in continuous_features:
    ax = plt.subplot(12,3,plotnumber)
    sns.boxplot(x="deposit", y= df[feature], data=df)
    plt.xlabel('deposit')
    plotnumber+=1
plt.show()
```



f) **Find Outliers in numerical features:**

Here I have plotted the outliers of the numerical features through boxplot.

```
3]: #boxplot on numerical features to find outliers
plt.figure(figsize=(20,60), facecolor='white')
plotnumber =1
for numerical_feature in numerical_features:
    ax = plt.subplot(12,3,plotnumber)
    sns.boxplot(df[numerical_feature])
    plt.xlabel(numerical_feature)
    plotnumber+=1
plt.show()
```



g) **Explore the Correlation between numerical features:**

Here I have plotted the correlation between the numerical features through heatmap

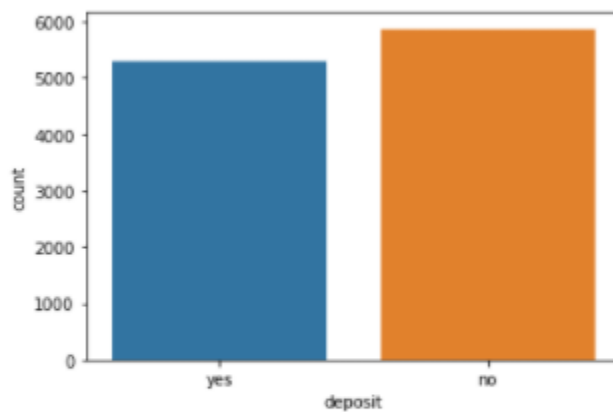


h) Check the Data set is balanced or not based on target values in classification:

```

In [25]: #total patient count based on cardio_results
sns.countplot(x='deposit',data=df)
plt.show()

```



```

In [26]: df['deposit'].groupby(df['deposit']).count()

```

```

Out[26]: deposit
no      5873
yes     5289
Name: deposit, dtype: int64

```

ii) Feature Engineering:

In this stage , I dropped unwanted features , handled the missing values,categorical features, feature scaling and also removed the outliers.

a) Drop unwanted Features:

```
[30]: #default features does not play imp role
df2.groupby(['deposit','default']).size()

In[30]: deposit    default
no         no      5757
         yes      116
yes        no     5237
         yes       52
dtype: int64

In[31]: df2.drop(['default'],axis=1, inplace=True)

In[32]: df2.groupby(['deposit','pdays']).size()

In[32]: deposit    pdays
no         -1      4940
         1         2
         2         6
         5         2
         6         2
         ...
yes        804         1
         805         1
         828         1
         842         1
         854         1
Length: 732, dtype: int64

In[33]: # drop pdays as it has -1 value for around 40%+
df2.drop(['pdays'],axis=1, inplace=True)
```

b) Handle Categorical Features

Handling categorical feature

```
In [42]: cat_columns = ['job', 'marital', 'education', 'contact', 'month', 'outcome']
for col in cat_columns:
    df4 = pd.concat([df4.drop(col, axis=1), pd.get_dummies(df4[col], prefix=col, prefix_sep='_', drop_first=True, dummy_na=False)], axis=1)
```

```
In [43]: bool_columns = ['housing', 'loan', 'deposit']
for col in bool_columns:
    df4[col+'_new'] = df4[col].apply(lambda x : 1 if x == 'yes' else 0)
    df4.drop(col, axis=1, inplace=True)
```

```
In [44]: df4.head()
```

```
Out[44]:
```

	age	balance	day	duration	campaign	previous	job_blue-collar	job_entrepreneur	job_housemaid	job_management	...	month_may	month_nov	month_oct
0	59	2343	5	1042	1	0	0	0	0	0	...	1	0	0
1	58	45	5	1467	1	0	0	0	0	0	...	1	0	0
2	41	1270	5	1389	1	0	0	0	0	0	...	1	0	0
3	55	2476	5	579	1	0	0	0	0	0	...	1	0	0
4	54	184	5	673	2	0	0	0	0	0	...	1	0	0

5 rows x 41 columns

c) Handle Feature Scaling:

Feature scaling of numerical features

```
n [ ]: numerical_features

n [ ]: feature_to_scale=[]
      for i in numerical_features:
          if i=='pdays':          # because 'pdays' was dropped
              continue
          else:
              feature_to_scale.append(i)

n [ ]: feature_to_scale

n [ ]: from sklearn.preprocessing import StandardScaler
      scaled_features = df4.copy()
      col_names = feature_to_scale
      features = scaled_features[col_names]
      scaler = StandardScaler().fit(features.values)
      features = scaler.transform(features.values)
      scaled_features[col_names] = features
      print(scaled_features)

n [ ]: scaled_features_df = pd.DataFrame(scaled_features, index=df4.index, columns=df4.columns)

n [ ]: scaled_features_df
```

iii) Train test split:

Split Dataset into Training set and Test set

```
[46]: X = df4.drop(['deposit_new'],axis=1)
      y = df4['deposit_new']
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.25)

[47]: len(X_train)

[47]: 8364

[48]: len(X_test)

[48]: 2788
```

iv) Feature Scaling:

Feature scaling


```
In [49]: from sklearn.preprocessing import StandardScaler  
sc = StandardScaler()  
sc.fit(X_train)  
X_train = sc.transform(X_train)  
  
X_test = sc.transform(X_test)
```

4. Modelling:

i) SVM:

Definition:

“Support Vector Machine” (SVM) is a supervised machine learning algorithm that can be used for both classification or regression challenges. However, it is mostly used in classification problems. In the SVM algorithm, we plot each data item as a point in n-dimensional space (where n is a number of features you have) with the value of each feature being the value of a particular coordinate. Then, we perform classification by finding the hyper-plane that differentiates the two classes very well.



Performance Metrics:

Accuracy: The accuracy of a [machine learning](#) classification algorithm is one way to measure how often the algorithm [classifies](#) a data point correctly. Accuracy is the number of correctly predicted data points out of all the data points.

Here in SVM model, Accuracy is 0.8238880918220947

Precision: Precision attempts to answer what proportion of positive identifications was actually correct. Precision is defined as follows:

$$precision = \frac{tp}{tp + fp}$$

Here in SVM model, Precision is 0.8262518968133535

Recall: Recall attempts to answer what proportion of actual positives was identified correctly. Mathematically, recall is defined as follows:

$$Recall = \frac{True\ Positive}{True\ Positive + False\ Negative}$$

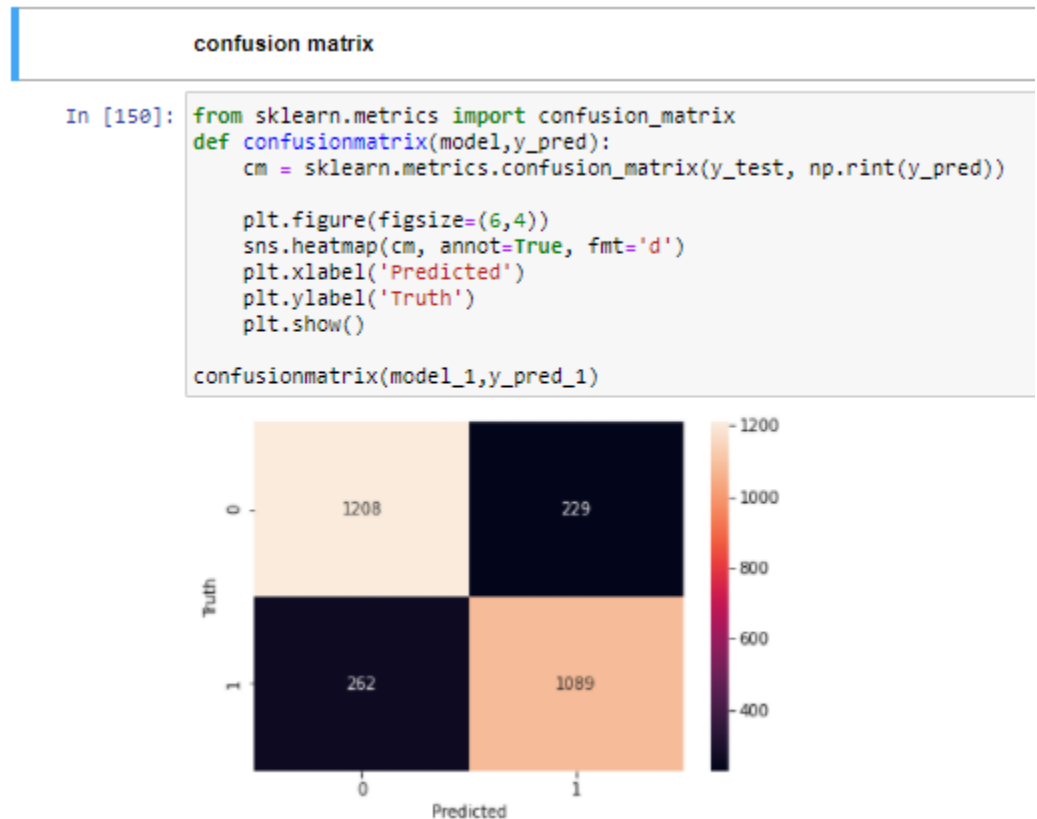
Here in SVM model, Recall is 0.8060695780903034

Confusion Matrix:

A Confusion matrix is an N x N matrix used for evaluating the performance of a classification model, where N is the number of target

classes. The matrix compares the actual target values with those predicted by the machine learning model.

Here in SVM Model confusion matrix is:




ROC Curve:

An **ROC curve (receiver operating characteristic curve)** is a graph showing the performance of a classification model at all classification thresholds. This curve plots two parameters:

- True Positive Rate
- False Positive Rate

Here in SVM model, ROC Curve is represented as:



Where area under curve (AUC) is .82 which is quite good.

Screenshot of the output:

```

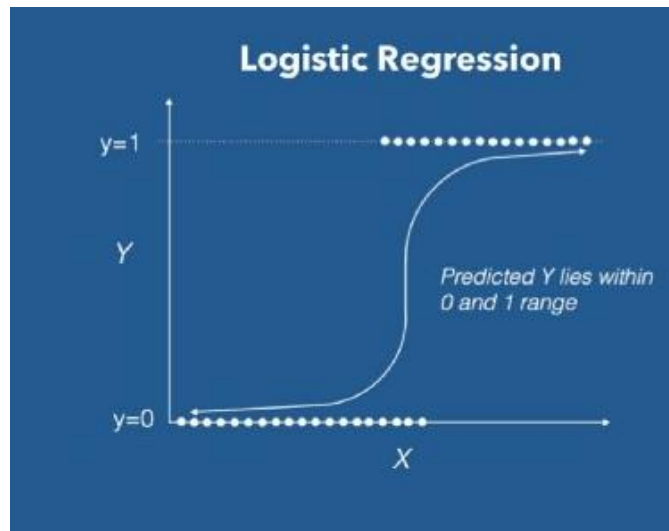
Accuracy: 0.8238880918220947
Precision: 0.8262518968133535
Recall: 0.8060695780903034

```

	precision	recall	f1-score	support
0	0.82	0.84	0.83	1437
1	0.83	0.81	0.82	1351
accuracy			0.82	2788
macro avg	0.82	0.82	0.82	2788
weighted avg	0.82	0.82	0.82	2788

ii) Logistic Regression:

Definition: Logistic Regression is a Machine Learning algorithm which is used for the classification problems, it is a predictive analysis algorithm and based on the concept of probability. We can call a Logistic Regression a Linear Regression model but the Logistic Regression uses a more complex cost function, this cost function can be defined as the '**Sigmoid function**' or also known as the 'logistic function' instead of a linear function.



Performance Metrics:


Accuracy : The accuracy of this model is: 0.833931133428

Precision: The precision of this model is: 0.8404255319

Recall: The recall of this model is: 0.8138337012

Confusion Matrix:

```
[151]: confusionmatrix(model_2,y_pred_2)
```




ROC Curve:

```
ROC_curve(y_pred_2)
```

```
[0. 0.49547669 1. ]  
[0. 0.44707624 1. ]  
[2 1 0]
```

<Figure size 432x288 with 0 Axes>



The AUC here is .48 which is a poor value

Screenshot of the output:

```
Accuracy: 0.8339311334289814
Precision: 0.8404255319148937
Recall: 0.8138337012509198
      precision    recall  f1-score   support

     0       0.83       0.85       0.84       1429
     1       0.84       0.81       0.83       1359

 accuracy          0.83          0.83          0.83          2788
  macro avg       0.83          0.83          0.83          2788
 weighted avg     0.83          0.83          0.83          2788
```

iii) K-nearest neighbors:

Definition:

The k-nearest neighbors (KNN) algorithm is a simple, supervised machine learning algorithm that can be used to solve both classification and regression problems. It's easy to implement and understand, but has a major drawback of becoming significantly slower as the size of that data in use grows.

KNN works by finding the distances between a query and all the examples in the data, selecting the specified number examples (K) closest to the query, then votes for the most frequent label (in the case of classification)

Performance Metrics:


Accuracy : The accuracy of this model is : 0.75322812

Precision: The precision of this model is : 0.81541389

Recall: The recall of this model is : 0.63434492

Confusion Matrix:

```
In [160]: confusionmatrix(model_3,y_pred_3)
```




ROC Curve:

```
ROC_curve(y_pred_3)
```

```
[0. 0.13500348 1.  ]  
[0. 0.63434493 1.  ]  
[2 1 0]
```

<Figure size 432x288 with 0 Axes>



AUC is .75 ,which is quite good.

Tuning Details:

i) `n_neighbors = 3:`

Accuracy: 0.7553802008608321

ii) n_neighbors = 21:

Accuracy: 0.75322812051

Screenshot of the output:

```
Accuracy: 0.7532281205164992
Precision: 0.8154138915318744
Recall: 0.6343449296817173
```

	precision	recall	f1-score	support
0	0.72	0.86	0.78	1437
1	0.82	0.63	0.71	1351
accuracy			0.75	2788
macro avg	0.77	0.75	0.75	2788
weighted avg	0.76	0.75	0.75	2788

iv) Decision Tree:

Definition:

Decision Tree is a Supervised learning technique that can be used for both classification and Regression problems, but mostly it is preferred for solving Classification problems. It is a tree-structured classifier, where internal nodes represent the features of a dataset, branches represent the decision rules and each leaf node represents the outcome. In a Decision tree, there are two nodes, which are the Decision Node and Leaf Node. Decision nodes are used to make any decision and have multiple branches, whereas Leaf nodes are the output of those decisions and do not contain any further branches. The decisions or the test are performed on the basis of features of the given dataset.

Performance Metrics:


Accuracy : The accuracy of this model is : 0.77690100

Precision: The precision of this model is : 0.7742663

Recall: The recall of this model is : 0.7616580

Confusion Matrix:

```
In [167]: confusionmatrix(model_4,y_pred_4)
```




ROC Curve:

```
178]: ROC_curve(y_pred_4)
```

```
[0.      0.20876827 1.      ]  
[0.      0.76165803 1.      ]  
[2 1 0]
```

<Figure size 432x288 with 0 Axes>



The AUC is .78 , which is a good value.

Screenshot of the output:

Accuracy: 0.7769010043041606

Precision: 0.7742663656884876

Recall: 0.7616580310880829


	precision	recall	f1-score	support
0	0.78	0.79	0.79	1437
1	0.77	0.76	0.77	1351
accuracy			0.78	2788
macro avg	0.78	0.78	0.78	2788
weighted avg	0.78	0.78	0.78	2788

v) Random Forest:

Definition:

Random forest, like its name implies, consists of a large number of individual decision trees that operate as an **ensemble**. Each individual tree in the random forest spits out a class prediction and the class with the most votes becomes our model's prediction. The fundamental concept behind random forest is a simple but powerful one — the wisdom of crowds. In data science speak, the reason that the random forest model works so well is:

A large number of relatively uncorrelated models (trees) operating as a committee will outperform any of the individual constituent models.



Performance Metrics:


Accuracy : The accuracy of this model is : 0.858680

Precision: The precision of this model is : 0.82841

Recall: The recall of this model is : 0.893412

Confusion Matrix:

```
In [171]: confusionmatrix(model_5,y_pred_5)
```




ROC Curve:

```
ROC_curve(y_pred_5)
```

```
[0.      0.17397356 1.      ]  
[0.      0.89341229 1.      ]  
[2 1 0]
```

<Figure size 432x288 with 0 Axes>



The AUC is .86 ,which is quite good value.

Screenshot of the output:

```
Accuracy: 0.8586800573888091
Precision: 0.8284145504461222
Recall: 0.8934122871946706
```

	precision	recall	f1-score	support
0	0.89	0.83	0.86	1437
1	0.83	0.89	0.86	1351
accuracy			0.86	2788
macro avg	0.86	0.86	0.86	2788
weighted avg	0.86	0.86	0.86	2788


vi) Neural Network:

Definition:

Artificial neural networks are relatively crude electronic networks of neurons based on the neural structure of the brain. They process records one at a time, and learn by comparing their classification of the record (i.e., largely arbitrary) with the known actual classification of the record. The errors from the initial classification of the first record is fed back into the network, and used to modify the networks algorithm for further iterations.


A neuron in an artificial neural network is

1. A set of input values (x_i) and associated weights (w_i).
2. A function (g) that sums the weights and maps the results to an output (y).



Neurons are organized into layers: input, hidden and output. The input layer is composed not of full neurons, but rather consists simply of the record's values that are inputs to the

next layer of neurons. The next layer is the hidden layer. Several hidden layers can exist in one neural network. The final layer is the output layer, where there is one node for each class. A single sweep forward through the network results in the assignment of a value to each output node, and the record is assigned to the class node with the highest value.



Performance Metrics:


Accuracy : The accuracy of this model is : 0.83

Precision: The precision of this model is : 0.83


Recall: The recall of this model is : 0.83

Confusion Matrix:

```
n [184]: confusionmatrix(model_6,y_pred_6)
```



ROC Curve:



The AUC is .84, which is a good value.

Screenshot of the output:

	precision	recall	f1-score	support
0	0.86	0.81	0.83	1437
1	0.81	0.86	0.83	1351
accuracy			0.83	2788
macro avg	0.83	0.84	0.83	2788
weighted avg	0.84	0.83	0.83	2788